72

```
1
2
                    | PCI Pass-through in Xen ARM |
3
                    -----
 4
                   manish.jaggi@caviumnetworks.com
5
                    ------
6
7
                           Draft-3
8
9
10
11
    Introduction
12
    This document describes the design for the PCI passthrough support in Xen ARM.
13
    The target system is an ARM 64bit Soc with GICv3 and SMMU v2 and PCIe devices.
14
15
16
17
    Revision History
18
    Changes from Draft-1:
19
20
    a) map mmio hypercall removed from earlier draft
21
22
    b) device bar mapping into guest not 1:1
    c) holes in guest address space 32bit / 64bit for MMIO virtual BARs
23
24
    d) xenstore device's BAR info addition.
25
26
    Changes from Draft-2:
    ______
27
28
    a) DomU boot information updated with boot-time device assignment and hotplug.
    b) SMMU description added
29
30
    c) Mapping between streamID - bdf - deviceID.
    d) assign_device hypercall to include virtual(guest) sbdf.
31
32
    Toolstack to generate guest sbdf rather than pciback.
33
34
35
    Index
36
37
    (1) Background
38
39
    (2) Basic PCI Support in Xen ARM
    (2.1) pci hostbridge and pci hostbridge ops
40
    (2.2) PHYSDEVOP_HOSTBRIDGE_ADD hypercall
41
42
    (3) SMMU programming
43
44
    (3.1) Additions for PCI Passthrough
45
    (3.2) Mapping between streamID - deviceID - pci sbdf
46
    (4) Assignment of PCI device
47
48
    (4.1) Dom0
49
    (4.1.1) Stage 2 Mapping of GITS_ITRANSLATER space (4k)
50
    (4.1.1.1) For Dom0
51
52
    (4.1.1.2) For DomU
    (4.1.1.2.1) Hypercall Details: XEN_DOMCTL_get_itranslater_space
53
54
55
    (4.2) DomU
56
    (4.2.1) Reserved Areas in guest memory space
    (4.2.2) New entries in xenstore for device BARs
57
    (4.2.4) Hypercall Modification for bdf mapping notification to xen
58
59
60
    (5) DomU FrontEnd Bus Changes
    (5.1) Change in Linux PCI FrontEnd - backend driver for MSI/X programming
61
    (5.2) Frontend bus and interrupt parent vITS
62
63
    (6) NUMA and PCI passthrough
64
65
66
67

    Background of PCI passthrough

68
69
    Passthrough refers to assigning a pci device to a guest domain (domU) such that
70
    the guest has full control over the device. The MMIO space and interrupts are
    managed by the guest itself, close to how a bare kernel manages a device.
71
```

```
73
     Device's access to guest address space needs to be isolated and protected. SMMU
     (System MMU - IOMMU in ARM) is programmed by xen hypervisor to allow device
 74
 75
     access guest memory for data transfer and sending MSI/X interrupts. PCI devices
     generated message signalled interrupt write are within guest address spaces which
 76
 77
     are also translated using SMMU.
     For this reason the GITS (ITS address space) Interrupt Translation Register
 78
 79
     space is mapped in the guest address space.
80
     2. Basic PCI Support for ARM
 81
 82
 83
     The apis to read write from pci configuration space are based on segment:bdf.
     How the sbdf is mapped to a physical address is under the realm of the pci
84
 85
     host controller.
 86
     ARM PCI support in Xen, introduces pci host controller similar to what exists
 87
88
     in Linux. Each drivers registers callbacks, which are invoked on matching the
 89
     compatible property in pci device tree node.
 90
91
     2.1 pci_hostbridge and pci_hostbridge_ops
92
     The init function in the pci host driver calls to register hostbridge callbacks:
 93
94
     int pci_hostbridge_register(pci_hostbridge_t *pcihb);
95
96
     struct pci_hostbridge_ops {
97
     u32 (*pci_conf_read)(struct pci_hostbridge*, u32 bus, u32 devfn,
98
     u32 reg, u32 bytes);
     void (*pci_conf_write)(struct pci_hostbridge*, u32 bus, u32 devfn,
99
100
     u32 reg, u32 bytes, u32 val);
101
102
103
     struct pci_hostbridge{
104
     u32 segno;
105
     paddr_t cfg_base;
     paddr_t cfg_size;
106
     struct dt_device_node *dt_node;
107
108
     struct pci_hostbridge_ops ops;
109
     struct list_head list;
110
111
112
     A pci conf read function would internally be as follows:
     u32 pcihb_conf_read(u32 seg, u32 bus, u32 devfn,u32 reg, u32 bytes)
113
114
115
     pci hostbridge t *pcihb;
     list_for_each_entry(pcihb, &pci_hostbridge_list, list)
116
117
118
     if(pcihb->segno == seg)
119
     return pcihb->ops.pci conf read(pcihb, bus, devfn, reg, bytes);
120
121
     return -1;
122
     }
123
124
     2.2 PHYSDEVOP pci host bridge add hypercall
125
126
     Xen code accesses PCI configuration space based on the sbdf received from the
127
     guest. The order in which the pci device tree node appear may not be the same
     order of device enumeration in dom0. Thus there needs to be a mechanism to bind
128
     the segment number assigned by dom0 to the pci host controller. The hypercall
129
     is introduced:
130
131
     #define PHYSDEVOP_pci_host_bridge_add 44
132
133
     struct physdev_pci_host_bridge_add {
     /* IN */
134
135
     uint16_t seg;
     uint64_t cfg_base;
136
     uint64_t cfg_size;
137
138
139
     This hypercall is invoked before dom0 invokes the PHYSDEVOP_pci_device_add
140
141
     hypercall. The handler code invokes to update segment number in pci_hostbridge:
142
143
     int pci_hostbridge_setup(uint32_t segno, uint64_t cfg_base, uint64_t cfg_size);
144
```

```
Subsequent calls to pci_conf_read/write are completed by the pci_hostbridge_ops
145
     of the respective pci hostbridge.
146
147
     2.3 Helper Functions
148
149
150
     a) pci hostbridge dt node(pdev->seg);
151
     Returns the device tree node pointer of the pci node from which the pdev got
152
     enumerated.
153
154
     3. SMMU programming
155
156
157
     3.1. Additions for PCI Passthrough
158
159
     3.1.1 - add_device in iommu_ops is implemented.
160
161
     This is called when PHYSDEVOP_pci_add_device is called from dom0.
162
163
     .add_device = arm_smmu_add_dom0_dev,
164
     static int arm_smmu_add_dom0_dev(u8 devfn, struct device *dev)
165
166
     if (dev_is_pci(dev)) {
167
     struct pci_dev *pdev = to_pci_dev(dev);
168
     return arm_smmu_assign_dev(pdev->domain, devfn, dev);
169
170
     return -1;
171
     }
172
173
     3.1.2 dev_get_dev_node is modified for pci devices.
174
     The function is modified to return the dt_node of the pci hostbridge from
175
176
     the device tree. This is required as non-dt devices need a way to find on
     which smmu they are attached.
177
178
179
     static struct arm_smmu_device *find_smmu_for_device(struct device *dev)
180
181
     struct device_node *dev_node = dev_get_dev_node(dev);
182
183
     static struct device_node *dev_get_dev_node(struct device *dev)
184
185
     if (dev_is_pci(dev)) {
186
187
     struct pci_dev *pdev = to_pci_dev(dev);
     return pci_hostbridge_dt_node(pdev->seg);
188
189
     }
190
191
192
193
     3.2. Mapping between streamID - deviceID - pci sbdf - requesterID
     ______
194
     For a simpler case all should be equal to BDF. But there are some devices that
195
196
     use the wrong requester ID for DMA transactions. Linux kernel has pci quirks
197
     for these. How the same be implemented in Xen or a diffrent approach has to be
198
     taken is TODO here.
199
     Till that time, for basic implementation it is assumed that all are equal to BDF.
200
201
202
     4. Assignment of PCI device
203
204
205
     4.1 Dom0
206
207
     All PCI devices are assigned to dom0 unless hidden by pci-hide bootargs in dom0.
     Dom0 enumerates the PCI devices. For each device the MMIO space has to be mapped
208
     in the Stage2 translation for dom0. For dom0 xen maps the ranges from dt pci
209
210
     nodes in stage 2 translation during boot.
211
212
     4.1.1 Stage 2 Mapping of GITS_ITRANSLATER space (64k)
213
214
     GITS_ITRANSLATER space (64k) must be programmed in Stage2 translation so that SMMU
215
     can translate MSI(x) from the device using the page table of the domain.
216
```

```
217
218
     4.1.1.1 For Dom0
219
     GITS ITRANSLATER address space is mapped 1:1 during dom0 boot. For dom0 this
220
221
     mapping is done in the vgic driver. For domU the mapping is done by toolstack.
222
223
     4.1.1.2 For DomU
224
225
     For domU, while creating the domain, the toolstack reads the IPA from the
226
     macro GITS ITRANSLATER SPACE from xen/include/public/arch-arm.h. The PA is
227
     read from a new hypercall which returns the PA of the GITS_ITRANSLATER_SPACE.
     Subsequently the toolstack sends a hypercall to create a stage 2 mapping.
228
229
230
     Hypercall Details: XEN DOMCTL get itranslater space
231
232
     /* XEN_DOMCTL_get_itranslater_space */
233
     struct xen_domctl_get_itranslater_space {
234
     /* OUT variables. */
     uint64_aligned_t start_addr;
235
     uint64_aligned_t size;
236
237
238
     typedef struct xen_domctl_get_itranslater_space xen_domctl_get_itranslater_space;
239
     DEFINE_XEN_GUEST_HANDLE(xen_domctl_get_itranslater_space;
240
241
     4.2 DomU
242
243
     There are two ways a device is assigned
244
     In the flow of pci-attach device, the toolstack will read the pci configuration
245
     space BAR registers. The toolstack has the guest memory map and the information
246
     of the MMIO holes.
247
     When the first pci device is assigned to domU, toolstack allocates a virtual
248
249
     BAR region from the MMIO hole area. toolstack then sends domctl
250
     xc_domain_memory_mapping to map in stage2 translation.
251
252
     4.2.1 Reserved Areas in guest memory space
253
254
     Parts of the guest address space is reserved for mapping assigned pci device's
255
     BAR regions. Toolstack is responsible for allocating ranges from this area and
256
     creating stage 2 mapping for the domain.
257
258
     /* For 32bit */
     GUEST_MMIO_BAR_BASE_32, GUEST_MMIO_BAR_SIZE_32
259
260
261
     /* For 64bit */
262
     GUEST MMIO BAR BASE 64, GUEST MMIO BAR SIZE 64
263
264
     Note: For 64bit systems, PCI BAR regions should be mapped from
265
     GUEST MMIO BAR BASE 64.
266
267
     IPA is allocated from the {GUEST MMIO BAR BASE 64, GUEST MMIO BAR SIZE 64}
268
     range and PA is the values read from the BAR registers.
269
270
     4.2.2 New entries in xenstore for device BARs
271
272
273
     toolstack also updates the xenstore information for the device
274
     (virtualbar:physical bar). This information is read by xenpciback and returned
275
     to the pcifront driver configuration space reads for BAR.
276
     Entries created are as follows:
277
     /local/domain/0/backend/pci/1/0
278
279
     vdev-N
     BDF = ""
280
     BAR-0-IPA = ""
281
     BAR-0-PA = ""
282
     BAR-0-SIZE = ""
283
284
     BAR-M-IPA = ""
285
     BAR-M-PA = ""
286
     BAR-M-SIZE = ""
287
288
```

318

319

320 321

322

device node

```
Note: Is BAR M SIZE is 0, it is not a valied entry.
289
290
291
     4.2.4 Hypercall Modification for bdf mapping notification to xen
292
293
     Guest devfn generation currently done by xen-pciback to be done by toolstack
294
     only. Guest devfn is generated at the time of domain creation (if pci devices
295
     are specified in cfg file) or using xl pci-attach call.
296
297
     DomU FrontEnd Bus Changes
298
299
     5.1 Change in Linux PCI ForntEnd - backend driver for MSI/X programming
300
     ______
301
     FrontEnd backend communication for MSI is removed in XEN ARM. It would be
302
     handled by the gic-its driver in guest kernel and trapped in xen.
303
304
     5.2 Frontend bus and interrupt parent vITS
305
306
     On the Pci frontend bus msi-parent gicv3-its is added. As there is a single
     virtual its for a domU, as there is only a single virtual pci bus in domU. This
307
     ensures that the config msi calls are handled by the gicv3 its driver in domU
308
     kernel and not utilising frontend-backend communication between dom0-domU.
309
310
311
     It is required to have a gicv3-its node in guest device tree.
312
313
     6. NUMA domU and vITS
314
     a) On NUMA systems domU still have a single its node.
315
     b) How can xen identify the ITS on which a device is connected.
316
317
     - Using segment number query using api which gives pci host controllers
```

struct dt\_device\_node\* pci\_hostbridge\_dt\_node(uint32\_t segno)

c) Query the interrupt parent of the pci device node to find out the its.